

MA453/553 – Classwork 3.

Profiling Cython Programs

Due: 10/02/2024

1. Download the file `classwork3.zip` from your Canvas course page. Save it in your work directory (`~/MA453/Classwork/cw3`) and unzip it. The file `particle_simulator.py` contains the original Python function `evolve()`

```
1  def evolve(self, tmax, dt=0.00001):
2      nsteps = int(tmax/dt)
3      for i in range(nsteps):
4          for p in self.particles:
5              norm = (p.x**2 + p.y**2)**0.5
6              v_x = (-p.y)/norm
7              v_y = p.x/norm
8              d_x = dt * p.ang_speed * v_x
9              d_y = dt * p.ang_speed * v_y
10             p.x += d_x
11             p.y += d_y
```

and its improvements: `evolve_python()`, `evolve_numpy()` and `evolve_cython()`. Run the code with each of these functions by commenting/uncommenting function calls in the lines 129-132 and 179-186 and make sure you get the same results (trajectories).

```
$ cd ~/MA453/Classwork/cw3
$ python setup.py build_ext --inplace
$ python particle_simulator.py
```

2. Measure the timings of the code with the pure Python, NumPy and Cython versions.

```
$ python -m timeit -s "from particle_simulator import benchmark" "benchmark(100, 'python')"
$ python -m timeit -s "from particle_simulator import benchmark" "benchmark(100, 'numpy')"
$ python -m timeit -s "from particle_simulator import benchmark" "benchmark(100, 'numexpr')"
$ python -m timeit -s "from particle_simulator import benchmark" "benchmark(100, 'cython')"
```

Which version gives the best performance?

3. Why aren't we getting the best out of Cython function? Give a look to the file `cevolve.pyx`. The cython module `cevolve` uses the untyped version. Comment lines 9-20 and uncomment lines 23-47, rebuild the module and time the `benchmark()` again.

```
1 #def c_evolve(r_i, ang_speed_i, dt, nsteps):
2 #     v_i = np.empty_like(r_i)
3 #     for i in range(nsteps):
4 #         norm_i = np.sqrt((r_i ** 2).sum(axis=1))
5 #         v_i = r_i[:, [1, 0]]
6 #         v_i[:, 0] *= -1
7 #         v_i /= norm_i[:, np.newaxis]
8 #         d_i = dt * ang_speed_i[:, np.newaxis] * v_i
9 #         r_i += d_i
```

```
$ python setup.py build_ext --inplace
$ python -m timeit -s "from particle_simulator import benchmark" "benchmark(100, 'cython')"
```

How much improvement did you get?

```

1 def c_evolve(double[:, :] r_i, double[:] ang_speed_i, double dt, int nsteps):
2     cdef int i, j, nparticles = r_i.shape[0]
3     cdef double norm, x, y, vx, vy, dx, dy, ang_speed
4     for i in range(nsteps):
5         for j in range(nparticles):
6             x = r_i[j, 0]
7             y = r_i[j, 1]
8             ang_speed = ang_speed_i[j]
9             norm = sqrt(x ** 2 + y ** 2)
10            vx = (-y)/norm
11            vy = x/norm
12            dx = dt * ang_speed * vx
13            dy = dt * ang_speed * vy
14            r_i[j, 0] += dx
15            r_i[j, 1] += dy

```

4. Let us profile the Cython module `cevolve.pyx` with the `annotated view` option.

`$ cython -a cevolve.pyx` (This generates an HTML file `cevolve.html`)

Open the file `cevolve.html` and check which lines do have more interpreter-related calls.

`$ firefox cevolve.html`

The white lines corresponds translated C code, you can click these lines to see the code.

5. In the line `v_y=x/norm`, Cython checks that computed norm is not zero, otherwise it raises a `ZeroDivisorError`, and in the line `r_i[j,0]`, Cython checks if the indexes are within the bounds of the array. Add the following two lines before the function `c_evolve()` in the file `cevolve.pyx` and measure the timing of `benchmark(1000, 'cython')`.

```

1 @cython.boundscheck(False)
2 @cython.cdivision(True)
3 def c_evolve(double[:, :] r_i, double[:] w_i, double dt, int nsteps):

```

`$ python setup.py build_ext --inplace`

`$ python -m timeit -s "from particle_simulator import benchmark" "benchmark(100, 'cython')"`

6. Record all the timings from 2., 3 and 5. in a file `timing.txt` and send me using the mail command:

`$ mail -s "ma453:cw3" 453 < timings.txt`